

MOSAIC: Calibrated Per-Cell Uncertainty for Unpaired Single-Cell Multi-Omics Integration in Disease Atlases

Anonymous Author(s)

Anonymous Institution

ANONYMOUS@ANONYMOUS.EDU

Abstract

Disease atlases built from single-cell multi-omics measurements are now the de-facto substrate for precision-medicine hypothesis generation: which regulatory state drives a tumor subtype, which immune population expands in severe COVID-19, which enhancer is primed in a neurodevelopmental disorder. In practice, almost all such atlases are unpaired across modalities — the Human Cell Atlas transcriptome and the ENCODE chromatin-accessibility compendium describe overlapping populations through different molecular lenses, but no individual cell was measured in both. Every existing integration method returns a point-estimate cell-to-cell alignment with no signal to distinguish a trustworthy match from an artifact. The downstream consequence is clinically relevant: wrong alignments silently corrupt joint cluster discovery, regulatory inference, and spatial deconvolution — the steps that translate an atlas into a therapeutic hypothesis. We introduce MOSAIC, a method that turns unpaired multi-omics alignment into a well-defined probabilistic problem and reports a *cluster-resolved alignment entropy* as a per-cell uncertainty measure. MOSAIC combines per-modality information-bottleneck VAEs with a cluster-centroid cross-modal prediction target, an orthogonal-Procrustes post-alignment, and entropic Sinkhorn optimal transport whose transport plan is marginalized over partner cluster labels. The marginalization is the key technical insight: cell-level row entropy on the Sinkhorn plan *anti-correlates* with alignment error ($\rho = -0.65$) because of the geometry of dense clusters; cluster-resolved entropy removes this confound. On three paired benchmarks (PBMC 10x Multiome, E18 mouse brain 5k Multiome, and 10x CITE-seq), MOSAIC outperforms SCOT and uniPort baselines on every alignment metric and discriminates wrong-cluster assignments with AUROC 0.88–0.95. In a leave-one-cluster-out missing-type experiment designed to mimic disease-specific cell populations absent from a healthy reference, cluster-resolved entropy achieves mean AUROC 0.96–0.995. A cross-tissue negative control (PBMC RNA $\times$ mouse-brain ATAC, zero shared cell types) yields $4.2\times$ higher mean entropy than within-dataset runs, confirming the signal correctly reports high uncertainty when nothing truly aligns. MOSAIC provides the calibrated filtering substrate that clinical multi-omics integration currently lacks.

1. Introduction

Modern translational research increasingly relies on single-cell multi-omics atlases as the backbone of disease characterization. The Human Cell Atlas (Regev et al., 2017), the Human Tumor Atlas Network (Rozenblatt-Rosen et al., 2020), the COVID-19 Cell Atlas (Stephenson et al., 2021), and the HuBMAP project (HuBMAP Consortium, 2019) all operate on the same premise: by cataloguing transcriptional, epigenomic, and proteomic

states of individual cells in health and disease, we obtain the substrate from which biomarkers, drug targets, and mechanistic hypotheses are drawn. A central practical obstacle to this premise is that the relevant measurements are almost never taken on the same cell. Single-cell RNA-seq (scRNA-seq) and single-cell ATAC-seq (scATAC-seq) atlases profile overlapping cell populations through different molecular lenses, but each individual cell is destroyed at measurement; paired multi-omics assays (10x Multiome, CITE-seq) remain expensive and cover a small fraction of existing atlas data. Unpaired integration — mapping a cell measured in modality A to a corresponding cell measured in modality B — has therefore become a standard preprocessing step for joint clustering, regulatory inference, cell-communication analysis, and downstream disease-association studies (Demetci et al., 2022; Cao and Gao, 2022; Cao et al., 2022; Hao et al., 2024).

A clinically consequential gap pervades this literature: no method reports per-cell uncertainty about the match it produces. Existing approaches — Seurat v5 bridge integration (Hao et al., 2024), GLUE (Cao and Gao, 2022), SCOT (Demetci et al., 2022), uniPort (Cao et al., 2022), MultiVI (Ashuach et al., 2023), LIGER (Welch et al., 2019), Harmony (Korsunsky et al., 2019) — all return a point estimate. Some expose a nearest-neighbor rank; none tell a researcher whether a particular cell has been matched reliably. The downstream consequence is clinically important. Incorrect alignments propagate silently into joint cluster labels, differential-accessibility tests, enhancer–gene linkage inferences, and spatial deconvolution outputs — the steps that produce the therapeutic hypotheses these atlases are built to generate. A candidate regulatory program that is in fact an artifact of mismatched cells can redirect a pharmaceutical program by years; in a disease atlas, a cell population that exists only in the diseased tissue can be silently “matched” to whatever healthy population happens to be nearest in a shared embedding, producing a spurious null finding precisely where novelty is most clinically valuable.

We introduce MOSAIC (Multi-Omic Stochastic Alignment via Information-theoretic Coupling), a method that reframes unpaired single-cell multi-omics integration as a probabilistic problem and reports a cluster-resolved alignment entropy as a per-cell uncertainty measure. MOSAIC composes three components. Each modality is encoded by its own information-bottleneck variational autoencoder (IB-VAE) whose latent is constrained to both reconstruct the input and predict a cluster-centroid summary of the partner modality; the cluster-centroid target is a critical design choice, because naively using the partner cell’s own feature vector causes catastrophic per-cell memorization while the cluster-centroid formulation forces the latent to retain cluster identity — a generalizing signal — rather than individual cell fingerprints. The two independently-trained IB-VAE latents live in different rotations of the same underlying space, so we fit an orthogonal Procrustes rotation on shared cluster centroids to bring them into a common frame; this post-hoc alignment is cheap, interpretable, and avoids the joint-optimization complexity of GLUE and MultiVI. Finally we solve an entropic optimal-transport problem between the two latent clouds and, rather than reporting per-row entropy of the transport plan — which we show is miscalibrated as a per-cell uncertainty signal because of the geometry of dense clusters — we marginalize each row over the partner modality’s cluster labels. The resulting cluster-level distribution has a natural entropy that captures cross-cluster uncertainty and predicts (AUROC 0.88–0.95) which cells have been assigned to the wrong cluster.

The core technical insight, phrased as a negative result, is that the obvious candidate uncertainty signal — row entropy of an entropic-OT transport plan — is not merely weak but directionally wrong. On PBMC 10k it yields Spearman $\rho = -0.65$ against true alignment error: cells in dense clusters have many valid within-cluster candidates (inflating entropy) but also a closer true partner (reducing error), so the two effects reverse the expected correlation. We believe this pathology generalizes to any OT-based integration method that exposes per-row entropy as a confidence score, and the marginalization cure is cheap and principled. On three paired benchmarks — 10x Multiome PBMC 10k (immune cells, 11,303 paired cells, 18 cell types), 10x Multiome E18 mouse brain 5k (neural development, 4,531 cells, 20 types), and 10x CITE-seq PBMC (RNA + 14 protein markers, 16 types) — MOSAIC outperforms SCOT, uniPort, and non-IB raw-feature baselines on every alignment metric, identifies wrong-cluster cells with AUROC 0.88–0.95, and — most importantly for clinical use — in a leave-one-cluster-out missing-type test (the computational analog of “this cell population exists in the disease atlas but not in the healthy reference”) flags absent-type cells with mean AUROC 0.96 on PBMC, 0.995 on Brain, and 0.946 on CITE-seq. A cross-tissue negative control (PBMC RNA $\times$ mouse-brain ATAC, no shared cell types) shows 4.2 $\times$ higher mean cluster entropy than within-dataset controls, confirming the signal correctly reports high uncertainty when nothing truly aligns.

Generalizable Insights about Machine Learning in the Context of Healthcare

- **Unpaired multi-omics integration is the modal setting for clinical atlases, not a methodological afterthought.** Paired assays are expensive; the data that actually populates disease atlases is overwhelmingly unpaired across modalities. A method that reports only a point-estimate alignment leaves a clinician or translational researcher unable to distinguish a trustworthy regulatory hypothesis from a silent artifact.
- **Calibrated per-cell uncertainty is the clinically actionable primitive.** An alignment with quantified uncertainty enables filtering (discard low-confidence matches before downstream analysis), triage (prioritize high-confidence matches for validation), and detection of disease-specific populations that exist in one atlas but not in the reference. Clinical utility is unlocked by the uncertainty signal, not the alignment itself.
- **Naive entropy on OT transport plans is miscalibrated in a characteristic way.** Row entropy of an entropic Sinkhorn plan anti-correlates with alignment error because dense clusters jointly inflate entropy and shrink the true-partner distance. The fix — marginalize over partner cluster labels before taking the entropy — is cheap, principled, and, we conjecture, applicable to any OT-based integration method that wishes to expose per-cell confidence.
- **Cross-modal prediction targets must be cluster-level, not cell-level, to generalize.** In our development, switching from per-cell partner features to cluster-centroid targets produced a 3.6 $\times$ improvement in label-transfer accuracy without any other change. Any architecture that couples two independently-encoded modalities

through a cross-modal prediction head is susceptible to per-cell memorization unless the target has cluster-level structure.

2. Related Work

The unpaired multi-omics integration literature is crowded and largely driven by biological-benchmarking criteria: SCOT (Demetci et al., 2022) uses Gromov–Wasserstein optimal transport on intra-modality distance matrices; uniPort (Cao et al., 2022) uses coupled autoencoders with a domain-adversarial regularizer; MultiVI (Ashuach et al., 2023) jointly trains a shared-latent VAE with modality-specific decoders; GLUE (Cao and Gao, 2022) uses a graph-linked unified embedding constrained by prior biological graphs (e.g., peak-to-gene linkages); LIGER (Welch et al., 2019) and Harmony (Korsunsky et al., 2019) use non-negative matrix factorization and iterative clustering; and Seurat v5 bridge integration (Hao et al., 2024) re-uses a small paired dataset as an anchoring bridge between two unpaired datasets. None of these methods expose a per-cell probabilistic uncertainty on the match itself, and this is where MOSAIC is complementary: the IB-VAE + Procrustes + Sinkhorn scaffolding is standard (though the cluster-centroid cross-modal target we introduce is novel), and the main contribution is the cluster-resolved entropy as a calibrated uncertainty signal. Adjacent literature on uncertainty quantification in single-cell analysis — scANVI (Xu et al., 2021) exposes posterior variances on latent embeddings, uncertainty quantification has been applied to atlas-level cell-type transfer (Engelmann et al., 2022) — does not translate into cell-to-cell match confidence across modalities; no prior method, to our knowledge, reports per-cell uncertainty on an unpaired multi-omics match. Optimal transport for single-cell data is well-developed in temporal and spatial settings (Waddington-OT (Schiebinger et al., 2019), moscot (Klein et al., 2025), PASTE (Zeira et al., 2022), SCOT itself), and the row entropy of a Sinkhorn plan is sometimes used qualitatively as a fuzziness indicator; its directional miscalibration as a per-cell uncertainty measure (Spearman $\rho = -0.65$ against true error) has not been reported, and we view the marginalization cure as a portable contribution for this family. Clinical applications of single-cell atlases in tumor immune microenvironment characterization (Zheng et al., 2021), COVID-19 severity stratification (Stephenson et al., 2021), inflammatory bowel disease subtyping (Smillie et al., 2019), and autoimmune cell-population mapping (Zhang et al., 2021) each integrate unpaired datasets across modalities and tissues, and each is therefore susceptible to silent alignment errors that a calibrated uncertainty signal would catch.

3. Methods

MOSAIC takes two single-cell modalities A (e.g., RNA) and B (e.g., ATAC) measured on disjoint cell populations — or paired cells whose correspondence is treated as unknown — and produces a joint latent space in which paired cells map close together and a per-cell cluster-resolved alignment uncertainty score calibrated against true alignment error. The pipeline has three stages, laid out below as continuous prose because their interfaces are tightly coupled; the joint latent is shown in Fig. 1.

Each modality’s encoder is a 3-layer MLP $d_x \rightarrow 512 \rightarrow 256 \rightarrow 64$ (with 2×64 output for $\mu, \log \sigma$ of the latent posterior), and the decoder mirrors the encoder; for RNA the decoder

produces the three parameters of a zero-inflated negative binomial distribution over raw counts with cell-library-size scaling (reconstruction loss is the ZINB negative log-likelihood), for ATAC it produces per-peak Bernoulli logits against binarized peak accessibility (binary cross-entropy), and for CITE-seq proteins we use CLR-normalized values with a Gaussian likelihood. A cross-modal prediction head $g_\phi : z \rightarrow \mathbb{R}^{d_y}$ is attached to the latent; its target $y_{\text{cross}}(i)$ for cell i is the centroid, in the partner modality’s 50-dim LSI/PCA space, of all cells that share i ’s cluster assignment. The full IB-VAE loss for modality m is

$$\mathcal{L}_m = \mathcal{L}_{\text{recon},m} + \beta \cdot \text{KL}(q(z|x) \parallel \mathcal{N}(0, I)) + \lambda_{\text{pred}} \cdot \|g_\phi(z) - y_{\text{cross}}\|_2^2, \quad (1)$$

with β annealed linearly over 30 epochs from 0 to its target and $\lambda_{\text{pred}} = 5$; training uses AdamW at lr 10^{-3} with a 10-epoch cosine warmup and 200 total epochs, and early stopping is disabled because the cross-modal loss is non-monotonic during the first half of training. We use $\beta = 0.001$ as the recommended default (an early $\beta = 0.01$ was defaulted based on a single-seed outlier and is retained in tables for traceability). The cluster-centroid design choice deserves particular emphasis: our original design used the partner cell’s own feature vector as the cross-modal target and failed catastrophically (train MSE 0.06, val MSE 1.2 — classic per-cell memorization), but switching to a cluster-centroid target, in which every cell of cluster k receives the same target (the mean of the partner modality’s cluster- k cells), closed the train/val gap (0.007/0.009) and produced a $3.6\times$ improvement in label transfer ($0.19 \rightarrow 0.68$) without any other change. This is the pivotal architectural choice; the IB-VAE backbone is otherwise standard.

After independent per-modality training, the two latent clouds live in different rotations of the same 64-dim space and must be brought into a common frame. We fit an orthogonal rotation $R \in \text{O}(64)$ minimizing $\sum_k \|Rz_k^B - \bar{z}_k^A\|_2^2$ where $\bar{z}_k^A, \bar{z}_k^B$ are cluster- k centroids in each modality, and we apply the rotation to the B latent. A similarity-Procrustes variant with isotropic scale reduced the centroid residual $1.75 \rightarrow 0.49$ but destroyed within-cluster geometry (LT collapsed $0.68 \rightarrow 0.19$), so rotation-only is what we report. On the post-Procrustes latents we then solve

$$\min_{P \in \Pi(\mu, \nu)} \langle C, P \rangle - \epsilon H(P), \quad (2)$$

where $C_{ij} = \|z_i^A - Rz_j^B\|^2 / \text{median}(C)$ is the median-normalized squared-Euclidean cost, $\epsilon = 0.05$, and $\Pi(\mu, \nu)$ is the transport polytope with uniform marginals; we use POT’s regular (non-log-domain) Sinkhorn (Flamary et al., 2021), because the log-domain variant OOMs at $11,303 \times 11,303$ on 16 GB RAM (`scipy.logsumexp` allocates multi-GB buffers), and for per-cell uncertainty analysis we restrict to a seeded 5000-cell subsample. For each source cell i , the transport-plan row P_i is a distribution over partner cells; raw row entropy $H(P_i)$ anti-correlates with alignment error, so we instead compute the partner-cluster marginal

$$p(c | i) = \frac{\sum_{j: \text{cluster}(j)=c} P_{ij}}{\sum_j P_{ij}}, \quad H_{\text{cluster}}(i) = -\frac{1}{\log K} \sum_c p(c | i) \log p(c | i), \quad (3)$$

normalized to $[0, 1]$ by the number of partner clusters K ; $H_{\text{cluster}}(i)$ captures cross-cluster uncertainty while ignoring within-cluster ambiguity, and the argmax over c is MOSAIC’s predicted partner cluster for cell i .

Under sub-Gaussian concentration of per-type latent vectors, exact Procrustes recovery of type centroids across modalities, minimum between-type separation Δ , and Sinkhorn regularization $\epsilon = \sigma^2/\log n$, two guarantees follow. For any cell i of a type k shared across both modalities, the cluster-resolved entropy decays exponentially in sample size: $\mathbb{E}[H_{\text{cluster}}(i)] = O((\log K)^{-1} \exp(-n\Delta^2/\sigma^2))$ (Theorem 1). For a type k^* present in A but absent in B , the cluster-resolved entropy is bounded away from zero: $\mathbb{E}[H_{\text{cluster}}(i)] \geq \log K_{\text{present}}/\log K - O(\sigma/\Delta)$ (Theorem 2). A corollary is that there exists a threshold τ such that thresholding on $H_{\text{cluster}}(i) < \tau$ correctly classifies shared-type vs. absent-type cells with probability $1 - \exp(-c'n)$. Proof sketches are in Appendix 6.1. Empirically, on PBMC 10k with $K = 18$ and $\Delta/\sigma \approx 3\text{--}5$, we observe argmax-correct cell entropy 0.12, argmax-wrong 0.49, and leave-one-out missing-type AUROC 0.96 — in line with the theorem’s qualitative predictions.

THEOREM 1 (SHARED-TYPE CLUSTER ENTROPY $\rightarrow 0$):

For any cell i of a type k shared across both modalities: $\mathbb{E}[H_{\text{cluster}}(i)] = O((\log K)^{-1} \exp(-n\Delta^2/\sigma^2))$.

THEOREM 2 (ABSENT-TYPE CLUSTER ENTROPY BOUNDED AWAY FROM 0):

For a type k^* present in A but absent in B : $\mathbb{E}[H_{\text{cluster}}(i)] \geq \log K_{\text{present}}/\log K - O(\sigma/\Delta)$, where K_{present} is the number of types remaining in B .

One methodological detail bounds the claims we can make and deserves explicit statement. The cross-modal target y_{cross} is a centroid over cells sharing a cluster assignment, and for our paired benchmark datasets those cluster assignments are obtained by Leiden-clustering RNA and propagating labels to ATAC via the paired ground truth (a row-copy; see `src/data/preprocess.py`). The training loop itself never accesses the cell-level pairing: we verified by inspection that `pair_idx` appears only in the post-training evaluation block, so cell-level paired information does not leak into IB-VAE optimization. However, because y_{cross} is precomputed from paired cluster labels, the training signal carries cluster-level paired information, and for a genuinely unpaired application y_{cross} must be constructed from (a) joint clustering of concatenated latents in an EM-like outer loop, (b) an external label source such as celltypist or Azimuth, or (c) known biological cell-type labels. The reported FOSCTTM and label-transfer numbers therefore reflect recovery of cell-level pairing given cluster-level correspondences, and the cross-tissue negative control in §5 tests the uncertainty signal in a setting where no shared cluster structure exists at all.

4. Cohort: Datasets

We evaluate on three publicly available benchmarks that collectively span immune, neural, and proteomic modalities: the 10x Multiome PBMC 10k dataset, which profiles peripheral blood mononuclear cells from a healthy donor simultaneously for RNA and ATAC and yields 11,303 paired cells across 18 Leiden clusters (T, B, NK, monocyte, DC, and platelet subtypes) after QC at ≥ 200 genes and ≥ 500 UMIs; the 10x Multiome E18 mouse brain 5k dataset, which profiles embryonic-day-18 mouse brain for RNA and ATAC and yields 4,531 paired cells across 20 Leiden clusters (neuronal progenitors, radial glia, differentiating neurons, and oligodendrocyte precursors); and the 10x CITE-seq 10k PBMC dataset, which jointly measures RNA (18,000+ genes) and 14 surface-protein ADT markers on $\sim 10,000$




Figure 1: MOSAIC aligned latent — PBMC 10k. Four-panel UMAP of the 64-dimensional IB-VAE latent after Procrustes rotation. *Left to right:* RNA-only colored by Leiden cluster; ATAC-only (same embedding); joint overlay (circles = RNA, triangles = ATAC); joint colored by cluster-resolved entropy H_{cluster} (viridis scale, low = certain). Cells of the same immune type co-localize across modalities; high-entropy cells (yellow/white) concentrate at cluster boundaries and sparse populations.

cells across 16 clusters. The paired ground truth is used only at evaluation time; the training loop does not read it. RNA counts are normalized to 10^4 total per cell, log1p-transformed, and reduced to the top-2000 highly variable genes (scanpy `flavor="seurat_v3"`) on raw counts; ATAC peak counts are TF-IDF-normalized, log1p-transformed, and reduced via 50-component LSI (first component dropped); CITE-seq protein counts are CLR-normalized and reduced via a 14-component LSI embedding; and each modality receives an independent Leiden clustering at resolution 1.0 on its own $k = 15$ neighborhood graph. All three benchmarks are publicly released demonstration datasets from 10x Genomics with no identifying patient information — this submission contains no private health data — and the pipeline is directly applicable to institutional disease atlases with appropriate IRB approval at deployment time.

5. Results

5.1. Alignment quality, baselines, and the uncertainty signal

Figs. 1 and 2 show UMAP projections of the concatenated post-Procrustes latent on PBMC and Brain: RNA and ATAC cells of each cell type co-cluster, and the global cell-type topology is preserved. Against the three baselines that we could reliably run — SCOT (our Gromov–Wasserstein reimplementation), raw PCA/LSI followed by Procrustes (an ablation of the IB-VAE), and uniPort — MOSAIC outperforms on every alignment metric on all three datasets (Table 1, Fig. 3). On Brain 5k the margin is dramatic: FOSCTTM 0.05 vs. 0.47 for SCOT (89% better), label transfer 0.96 vs. 0.13 (7.2 $\times$), ARI 0.87 vs. 0.03 (35 $\times$). SCOT’s Gromov–Wasserstein solver fails to converge to a meaningful coupling on Brain and CITE-seq, and uniPort gives worse-than-random FOSCTTM (> 0.5) on all three datasets, which we attribute to uniPort’s diagonal-integration assumption of feature-space commonality that does not hold for RNA+ATAC (different feature types) or RNA+protein (different measurement scales). The $\lambda_{\text{pred}} = 0$ ablation is instructive: removing the cross-




Figure 2: MOSAIC aligned latent — Brain 5k. Same four-panel layout as Fig. 1 for E18 mouse brain. Neural progenitor, neuronal, and glial cell-type topology is preserved after Procrustes alignment. Entropy is concentrated at transitional states (radial glia→neuron boundaries) and rare clusters.

Table 1: Alignment quality vs. baselines. Bold = best per column per dataset. MOSAIC runs at $\beta = 0.001$ on Brain and CITE-seq, $\beta = 0.01$ on PBMC (reported single-seed for direct comparison to baselines; multi-seed results in Table 2).

Dataset	Method	FOSCTTM ↓	LT A→B ↑	LT B→A ↑	ARI ↑	Per-cell UQ?
PBMC 10k	MOSAIC ($\beta = 0.01$)	0.194	0.664	0.694	0.651	Yes
	SCOT	0.248	0.324	0.383	0.322	No
	Raw PCA/LSI + Procrustes	0.328	0.132	0.653	0.093	No
	uniPort	0.563	0.134	0.136	0.067	No
Brain 5k	MOSAIC ($\beta = 0.001$)	0.052	0.960	0.967	0.870	Yes
	SCOT	0.475	0.134	0.067	0.025	No
	Raw PCA/LSI + Procrustes	0.334	0.095	0.429	0.063	No
	uniPort	0.509	0.095	0.059	0.050	No
CITE-seq	MOSAIC ($\beta = 0.001$)	0.094	0.872	0.960	0.791	Yes
	SCOT	0.550	0.061	0.080	0.304	No
	Raw PCA/LSI + Procrustes	0.237	0.327	0.550	0.377	No
	uniPort	0.463	0.144	0.149	0.394	No

modal head yields FOSCTTM 0.378, LT 0.05/0.39, and ARI 0.026 on PBMC — essentially random alignment — confirming that the cross-modal head is the component that forces the IB-VAE latent to retain cross-modally-predictive information. Counterintuitively, this variant produces a *higher* AUROC (0.962 vs. 0.894) for H_{cluster} -based wrong-cell detection, but only because 75% of cells are wrong-assigned and any weakly-informative signal achieves high AUROC by sheer base-rate advantage, which is why AUROC must be reported alongside argmax accuracy throughout the paper: a high AUROC against a high base error rate is a much easier test than a high AUROC against a low one.

Table 2 reports 10-seed (PBMC $\beta = 0.001$) and 3-seed (all others) means $\pm$ std. The recommended default $\beta = 0.001$ is strictly or tied-better than $\beta = 0.01$ on Brain across every metric and on PBMC across every metric except argmax cluster accuracy (96.4% vs. 98.4%); PBMC $\beta = 0.001$ ARI has higher seed variance (std 0.103 vs. 0.009) because a weaker bottleneck retains finer per-cell detail, making downstream KMeans more sensitive




Figure 3: Baseline comparison on PBMC and Brain. MOSAIC outperforms SCOT, raw-feature Procrustes, and uniPort across all primary alignment metrics on both datasets.

to initialization on PBMC’s heavily imbalanced cluster-size distribution, while Brain’s more uniform cluster sizes make the weaker bottleneck uniformly beneficial. The central per-cell uncertainty result appears in Figs. 4 and 5: cell-level row entropy on the Sinkhorn plan has Spearman $\rho = -0.648$ (PBMC) and -0.52 (Brain) against true alignment error — the opposite of what a calibrated uncertainty should show — because cells in dense clusters have many valid within-cluster candidates (inflating entropy) and simultaneously a closer true partner (reducing error). Marginalizing over partner cluster labels removes this confound: MOSAIC’s argmax cluster accuracy is 98.4% on PBMC ($\beta = 0.01$) and 98.2% on Brain ($\beta = 0.001$), and H_{cluster} discriminates wrong-cluster cells with AUROC 0.883 ± 0.010 (PBMC) and 0.946 ± 0.017 (Brain). This is the calibrated per-cell uncertainty signal the paper claims, and it is structurally downstream of the cluster-centroid cross-modal target discussed in §3.

5.2. Clinically realistic stress tests: missing populations, cross-tissue control, calibration, and stability

A clinically realistic deployment pattern is to align a disease atlas against a healthy reference and ask which cells are of a disease-specific type absent from the reference. We simulate

Table 2: Multi-seed means $\pm$ std (3 seeds for Brain and CITE-seq; 10 seeds for PBMC $\beta = 0.001$). Best row per dataset in bold.

Dataset	β	FOSCTTM	LT A $\rightarrow$ B	LT B $\rightarrow$ A	ARI
PBMC 10k	0.01	0.188 $\pm$ 0.006	0.689 $\pm$ 0.005	0.771 $\pm$ 0.029	0.687 $\pm$ 0.009
PBMC 10k	0.001	0.118 $\pm$ 0.008	0.913 $\pm$ 0.074	0.909 $\pm$ 0.035	0.724 $\pm$ 0.103
Brain 5k	0.01	0.129 $\pm$ 0.004	0.877 $\pm$ 0.005	0.740 $\pm$ 0.025	0.408 $\pm$ 0.039
Brain 5k	0.001	0.049 $\pm$ 0.002	0.962 $\pm$ 0.002	0.966 $\pm$ 0.010	0.935 $\pm$ 0.003




Figure 4: PBMC 10k entropy miscalibration. (A) Cell-level row entropy H_{cell} vs. true alignment error: anti-correlates ($\rho = -0.65$, red trend line), the opposite of a useful uncertainty signal. (B) Cluster-resolved H_{cluster} vs. true error: positive correlation (blue line) — the marginalization cure. (C) Spatial UMAP heatmap of H_{cluster} : high uncertainty (yellow) concentrates at cluster boundaries and rare populations, consistent with biological interpretation.




Figure 5: Brain 5k entropy miscalibration. Same three-panel layout as Fig. 4. Row entropy has the wrong sign ($\rho \approx -0.5$); cluster-resolved entropy is correctly ordered. The UMAP heatmap reveals that transitional progenitor states and small rare clusters carry the highest uncertainty.

this with a leave-one-cluster-out experiment: for each candidate cluster, we remove its cells from modality B entirely, run Sinkhorn, and use H_{cluster} as the score for “is this source cell of the removed type?” The resulting AUROCs (Fig. 6, Table 3) are substantial across all three datasets; at $\beta = 0.001$ on Brain, every one of 18 clusters is detected with AUROC > 0.988 — near-perfect missing-type identification — while on CITE-seq 15/16 clusters are




Figure 6: Leave-one-cluster-out missing-type AUROC. For each cluster, we remove its cells from the partner modality and compute AUROC for “is this source cell of the removed type?” using H_{cluster} as the score — the computational analog of “cell population exists in the disease atlas but not in the healthy reference.”

Table 3: Missing-type detection AUROC. Clusters of size 30–1500 tested per dataset.

Dataset	β	Mean AUROC	Median AUROC	Range
PBMC 10k	0.01	0.960	0.987	0.671 – 1.000
Brain 5k	0.01	0.959	0.964	0.915 – 0.997
Brain 5k	0.001	0.995	0.995	0.988 – 1.000
CITE-seq	0.001	0.946	0.998	0.318 – 1.000

detected at $\text{AUROC} > 0.95$ and the single hard case (AUROC 0.318) is a cluster transcriptionally near-redundant with another cluster remaining in the target, a pattern consistent across all three datasets. For clinical deployment this pattern suggests a natural secondary check: if H_{cluster} is low but a cell population is suspected to be absent, verify cluster non-redundancy in the reference atlas. Pushing this further, we ran a cross-tissue negative control in which the trained PBMC RNA IB-VAE embedding is aligned against the trained mouse-brain ATAC embedding — PBMC cells are T/B/NK/monocyte/DC/platelet, brain cells are neurons/glia/oligodendrocyte precursors, and no shared cell types exist, so a calibrated UQ signal should report uniformly high uncertainty. The observed mean H_{cluster} is 0.635 ± 0.133 (cross-tissue) vs. 0.153 (within-PBMC) and 0.083 (within-Brain), a $4.2\times$ increase, with the cross-tissue distribution bimodal at ~ 0.55 and 0.75 (Fig. 7): the signal correctly flags every cross-tissue cell as highly uncertain.

The final piece of the uncertainty picture concerns the distinction between discrimination and calibration, and between single-seed and ensemble reliability. H_{cluster} discriminates correct-vs-wrong cell assignments with high AUROC but this is not the same as being calibrated in the strict sense of matching a probability, and the calibration curves in Fig. 8 make this explicit: binning cells by H_{cluster} into deciles and plotting mean H_{cluster} vs. observed error rate, the curves are monotonically increasing on all datasets, expected calibration error (ECE) ranges 0.05–0.16 with PBMC $\beta = 0.001$ best at 0.050, and Brier scores are




Figure 7: Cross-tissue negative control: PBMC RNA aligned against mouse-brain ATAC. No shared cell types exist. Mean cluster entropy is $4.2\times$ higher than within-dataset controls.




Figure 8: Calibration curves: mean H_{cluster} per decile vs. observed error rate. Monotonically increasing on all datasets; ECE 0.05–0.16.

0.024–0.052. The accurate characterization is therefore “well-discriminating per-cell uncertainty with monotonic risk ordering” rather than a calibrated probability; for deployment uses that require true probabilities, a post-hoc recalibration (Platt or isotonic) on a held-out set is recommended. Cross-seed reproducibility is moderate but not perfect: pairwise Spearman ρ of per-cell H_{cluster} across 10 PBMC $\beta = 0.001$ seeds is 0.644 ± 0.073 , top-10% flagged-cell Jaccard overlap is 0.172 vs. 0.053 random ($3.2\times$), so for high-stakes clinical filtering we recommend ensembling 3–5 seeds rather than trusting a single run.

6. Required: Discussion

MOSAIC is the first method we are aware of that provides per-cell uncertainty for unpaired single-cell multi-omics alignment, and the clinical utility is in the filtering, triage, and missing-type detection use cases that an uncaptured alignment error silently corrupts. The




Figure 9: Disease simulation: detecting clinically absent cell types. *Left:* Five PBMC immune disease scenarios (mean AUROC 0.952). *Right:* Five brain neurological disease scenarios (mean AUROC 0.993). Bar color encodes cluster-resolved entropy ratio (absent/present H_{cluster} , warmer = stronger signal). Dotted line = chance (AUROC 0.5); solid colored line = dataset mean.

technical contribution is the cluster-resolved marginalization: row entropy of an entropic Sinkhorn plan is directionally miscalibrated as a per-cell uncertainty signal because of the geometry of dense clusters, and marginalizing over partner cluster labels is a cheap, principled fix that we conjecture applies broadly to OT-based integration methods. Practitioner-facing deployment guidance falls out of the experiments above: for disease-atlas $\times$ healthy-reference comparison, deploy MOSAIC with $\beta = 0.001$ and flag all cells with H_{cluster} above the 75th percentile of a within-dataset control run on the healthy reference as candidate disease-specific populations warranting follow-up validation; for regulatory inference (e.g., enhancer-gene linkage from an unpaired RNA+ATAC integration), discard cells with H_{cluster} above a user-defined threshold before computing linkage statistics, as the clinical analog of standard quality-control filtering; for high-stakes filtering, ensemble 3–5 seeds because single-seed per-cell H_{cluster} has moderate but imperfect cross-seed stability.

6.1. Clinical and translational applications

To illustrate how H_{cluster} maps to specific disease contexts, we ran three targeted downstream analyses on the existing embeddings without any additional training. First, a *disease simulation* study: for five PBMC immune disease scenarios (CD8 T lymphopenia, NK deficiency, B cell aplasia, monocytopenia, Treg deficiency) and five brain neurological scenarios (excitatory neuron loss modelling Alzheimer’s disease, inhibitory neuron loss modelling epilepsy, oligodendrocyte loss modelling multiple sclerosis, astrocyte depletion, microglia depletion via PLX5622 model), we removed the target cell population from the ATAC reference and measured AUROC for detecting absent cells via H_{cluster} (Fig. 9). Mean AUROC was 0.952 on PBMC immune diseases and 0.993 on brain neurological conditions, with entropy ratios (absent/present mean H_{cluster}) of 3.2–4.8 $\times$ (immune) and 10–27 $\times$ (neural). The elevated ratio in brain is expected: E18 mouse brain clusters are more transcriptomically distinct than PBMC immune clusters, making absent-type signal sharper.




Figure 10: CITE-seq protein marker uncertainty analysis. (A) Mean protein expression difference (high-uncertainty minus low-uncertainty cells) for all 14 ADT markers. (B) Statistical significance ($-\log_{10}$ Mann-Whitney p -value). Red = enriched in uncertain cells (T/memory T markers); blue = depleted (NK/naive T markers). All markers significant at $p < 10^{-13}$.

Second, a *protein marker UQ* analysis on CITE-seq: we compared protein expression between the top and bottom 10% of cells ranked by H_{cluster} (Fig. 10). High-uncertainty cells are enriched for T cell markers (CD3, CD4, CD45RO, CD8a) and depleted for lineage-committed markers (CD45RA for naive T, CD56/CD16 for NK). This is biologically coherent: T cells, especially memory subtypes, exhibit the most transcriptome-to-proteome discordance because T cell activation induces rapid surface protein remodeling that lags behind transcriptional changes. The signal correctly identifies which cells are hardest to align and why.

Third, an *immune checkpoint immunotherapy* analysis: PD-1⁺TIGIT⁺ double-positive exhausted T cells (the primary target of checkpoint inhibitor therapy) show mean $H_{\text{cluster}} = 0.238$ vs. 0.201 for PD-1⁻TIGIT⁻ fresh cells ($p = 9.2 \times 10^{-5}$, Mann-Whitney; Fig. 11). CD4⁺PD-1⁺ Treg/exhausted-CD4 T cells have substantially higher entropy than CD8a⁺PD-1⁺ effector-exhausted T cells ($H = 0.241$ vs. 0.188, $p = 3.6 \times 10^{-18}$), suggesting that CD4 exhausted T cells have more uncertain transcriptome-to-proteome correspondence — clinically important because CD4 exhaustion status is a predictor of checkpoint inhibitor response, and alignment uncertainty here flags cells where transcriptomic readout may not reliably predict surface PD-1 expression.

The limitations are substantial and we state them plainly. The Procrustes rotation and H_{cluster} both require cluster labels, and paired-benchmark evaluation uses Leiden labels propagated via ground-truth pairing; in genuinely unpaired clinical deployment, labels must come from joint clustering, external references such as celltypist or Azimuth, or known cell-type annotations, and the downstream clinical interpretation of H_{cluster} is conditional on having a sensible cluster assignment upstream — poor cluster assignments will propagate into poor uncertainty estimates. MOSAIC as written handles two modalities, and generalization to three or more modalities requires a group Procrustes problem and a multi-modal cluster-marginal definition that we defer to future work. H_{cluster} is well-




Figure 11: Immune checkpoint immunotherapy entropy analysis (CITE-seq). Violin plots of H_{cluster} for checkpoint-high vs. checkpoint-low cell subsets. Exhausted T cells (PD-1⁺, TIGIT⁺, PD-1⁺TIGIT⁺ double-positive) show significantly higher alignment uncertainty, and CD4⁺PD-1⁺ cells are more uncertain than CD8a⁺PD-1⁺ cells, identifying the subset where transcriptome-proteome discordance is highest for immunotherapy biomarker readout.

discriminating (AUROC 0.88–0.95) but not strictly probability-calibrated (ECE 0.05–0.16), so downstream uses that require true probabilities need post-hoc recalibration on held-out data. The convergence theorem assumes Procrustes recovery, and finite-sample Procrustes residual contributes an additional $O(\text{residual}/\sigma)$ term to both the shared-type concentration and the absent-type bound, which is consistent with the empirical gap between the theorem’s asymptotic rates and the observed AUROC 0.96. Finally, we evaluate on three public demonstration datasets (PBMC Multiome, mouse brain Multiome, CITE-seq); clinical validation on a disease-context atlas with accompanying healthy reference is a natural next step and is in progress. Open questions for follow-up include end-to-end learning of the cluster assignment that H_{cluster} depends on (closing the loop on the label-free setting), multi-modal (≥ 3) generalization via group Procrustes, and a principled uncertainty decomposition for populations present in some but not all modalities (richer than the binary present/absent framing).

References

Tal Ashuach, Mariano I Gabitto, Rohan V Koodli, Giuseppe-Antonio Saldi, Michael I Jordan, and Nir Yosef. MultiVI: deep generative model for the integration of multimodal data. *Nature Methods*, 20(8):1222–1231, 2023.

- Kai Cao, Qiyu Gong, Yiguang Hong, and Lin Wan. A unified computational framework for single-cell data integration with optimal transport. *Nature Communications*, 13(1):7419, 2022.
- Zhi-Jie Cao and Ge Gao. Multi-omics single-cell data integration and regulatory inference with graph-linked embedding. *Nature Biotechnology*, 40(10):1458–1466, 2022.
- Pinar Demetci, Rebecca Santorella, Björn Sandstede, William Stafford Noble, and Ritambhara Singh. SCOT: Single-cell multi-omics alignment with optimal transport. *Journal of Computational Biology*, 29(1):3–18, 2022.
- Jan Engelman, Leon Hetzel, Giovanni Palla, Lisa Sikkema, Malte Luecken, and Fabian J Theis. Uncertainty quantification for atlas-level cell type transfer, 2022.
- Rémi Flamary, Nicolas Courty, Alexandre Gramfort, Mokhtar Z Alaya, Aurélie Boissunon, Stanislas Chambon, Laetitia Chapel, Adrien Corenflos, Kilian Fatras, Nemo Fournier, et al. POT: Python optimal transport. *Journal of Machine Learning Research*, 22(78):1–8, 2021.
- Yuhan Hao, Tim Stuart, Madeline H Kowalski, Saket Choudhary, Paul Hoffman, Austin Hartman, Avi Srivastava, Gesmira Molla, Shaista Madad, Carlos Fernandez-Granda, and Rahul Satija. Dictionary learning for integrative, multimodal and scalable single-cell analysis. *Nature Biotechnology*, 42:293–304, 2024.
- HuBMAP Consortium. The human body at cellular resolution: the nih human biomolecular atlas program. *Nature*, 574(7777):187–192, 2019.
- Dominik Klein, Giovanni Palla, Marius Lange, Michal Klein, Zoe Piran, Manuel Gander, Laetitia Meng-Papaxanthos, Michael Sterr, Lama Saber, Changying Jing, et al. Mapping cells through time and space with moscot. *Nature*, 638(8052):1065–1075, 2025.
- Ilya Korsunsky, Nghia Millard, Jean Fan, Kamil Slowikowski, Fan Zhang, Kevin Wei, Yuriy Baglaenko, Michael Brenner, Po-ru Loh, and Soumya Raychaudhuri. Fast, sensitive and accurate integration of single-cell data with harmony. *Nature Methods*, 16(12):1289–1296, 2019.
- Aviv Regev, Sarah A Teichmann, Eric S Lander, Ido Amit, Christophe Benoist, Ewan Birney, Bernd Bodenmiller, Peter Campbell, Piero Carninci, Menna Clatworthy, et al. The human cell atlas. *eLife*, 6:e27041, 2017.
- Orit Rozenblatt-Rosen, Aviv Regev, Philipp Oberdoerffer, Tal Nawy, Anna Hupalowska, Jennifer E Rood, Orr Ashenberg, Ethan Cerami, Robert J Coffey, Emek Demir, et al. The human tumor atlas network: charting tumor transitions across space and time at single-cell resolution. *Cell*, 181(2):236–249, 2020.
- Geoffrey Schiebinger, Jian Shu, Marcin Tabaka, Brian Cleary, Vidya Subramanian, Aryeh Solomon, Joshua Gould, Siyan Liu, Stacie Lin, Peter Berube, et al. Optimal-transport analysis of single-cell gene expression identifies developmental trajectories in reprogramming. *Cell*, 176(4):928–943, 2019.

- Christopher S Smillie, Moshe Biton, Jose Ordovas-Montanes, Keri M Sullivan, Grace Burgin, Daniel B Graham, Rebecca H Herbst, Noga Rogel, Michal Slyper, Julia Waldman, et al. Intra- and inter-cellular rewiring of the human colon during ulcerative colitis. *Cell*, 178(3):714–730, 2019.
- Emily Stephenson, Gary Reynolds, Rachel A Botting, Fernando J Calero-Nieto, Michael D Morgan, Zewen Kelvin Tuong, Karsten Bach, Waradon Sungnak, Kaylee B Worlock, Masahiro Yoshida, et al. Single-cell multi-omics analysis of the immune response in covid-19. *Nature Medicine*, 27(5):904–916, 2021.
- Joshua D Welch, Velina Kozareva, Ashley Ferreira, Charles Vanderburg, Carly Martin, and Evan Z Macosko. Single-cell multi-omic integration compares and contrasts features of brain cell identity. *Cell*, 177(7):1873–1887, 2019.
- Chenling Xu, Romain Lopez, Edouard Mehlman, Jeffrey Regier, Michael I Jordan, and Nir Yosef. Probabilistic harmonization and annotation of single-cell transcriptomics data with deep generative models. *Molecular Systems Biology*, 17(1):e9620, 2021.
- Ron Zeira, Max Land, Alexander Strzalkowski, and Benjamin J Raphael. Alignment and integration of spatial transcriptomics data. *Nature Methods*, 19(5):567–575, 2022.
- Fan Zhang, Joseph R Mears, Lorien Shakib, Jessica I Beynor, Sara Shanaj, Ilya Korsunsky, Aparna Nathan, Laura T Donlin, and Soumya Raychaudhuri. IFN- γ and TNF- α drive a CXCL10+CCL2+ macrophage phenotype expanded in severe COVID-19 and other diseases with tissue inflammation. *Genome Medicine*, 13(1):64, 2021.
- Liangtao Zheng, Shishang Qin, Wen Si, Anqiang Wang, Baocai Xing, Ranran Gao, Xianwen Ren, Li Wang, Xiaojuan Wu, Ji Zhang, et al. Pan-cancer single-cell landscape of tumor-infiltrating t cells. *Science*, 374(6574):abe6474, 2021.

Appendix A. Proof sketches, β sensitivity, compute, and reproducibility

We carry four assumptions throughout the theorem statements in §3: (A1) sub-Gaussian latent concentration, i.e., for each type k , $\mathbb{E}[\exp(\langle u, z_i^A - \mu_k^A \rangle)] \leq \exp(\sigma^2 \|u\|^2/2)$ and symmetrically for B ; (A2) exact Procrustes recovery, $\mu_k^A = \mu_k^B$ after rotation; (A3) between-type separation $\Delta = \min_{k \neq \ell} \|\mu_k - \mu_\ell\| > 0$; (A4) Sinkhorn regularization $\epsilon = \sigma^2/\log n$. For Theorem 1, by (A1) the type- k cluster in B forms a ball of radius $O(\sigma\sqrt{\log n})$ around μ_k^B with probability $1 - 1/n$, so for a source cell i of type k the distance to μ_k^B is $\leq 2\sigma\sqrt{\log n}$ and the distance to any $\mu_\ell^B, \ell \neq k$ is $\geq \Delta - O(\sigma\sqrt{\log n})$; the Sinkhorn log-sum-exp then gives $p(k | i) \geq \exp(-2\sigma^2 \log n/\epsilon)/[\exp(-2\sigma^2 \log n/\epsilon) + (K - 1)\exp(-\Delta^2/(2\epsilon) + O(\sigma^2 \log n/\epsilon))]$, and with $\epsilon = \sigma^2/\log n$ (A4) the numerator is $\exp(-2(\log n)^2)$ while the subdominant denominator term is $(K - 1)\exp(-\Delta^2 \log n/(2\sigma^2))$, so for large n , $p(k | i) \rightarrow 1$ and $H_{\text{cluster}}(i) \rightarrow 0$ at the stated rate. For Theorem 2, type k^* absent in B has its closest centroid at distance $\geq \Delta$; each remaining type k' receives mass $\sim \exp(-\Delta^2/(2\epsilon))$ but the $K - 1$ remaining types have near-equal nearest distances up to $O(\sigma/\Delta)$ corrections, so the cluster marginal is nearly uniform over K_{present} remaining types, yielding $H_{\text{cluster}}(i) \geq \log K_{\text{present}}/\log K - O(\sigma/\Delta)$. The corollary follows: choose $\tau \in (O(\log K^{-1} \exp(-n\Delta^2/\sigma^2)), \log K_{\text{present}}/\log K - O(\sigma/\Delta))$, a window that is nonempty for n polynomial in K and $1/\sigma$, and thresholding on $H_{\text{cluster}}(i) < \tau$ classifies shared-type cells (low entropy) and absent-type cells (high entropy) with exponentially small error.

Fig. 12 shows β sensitivity on PBMC and Brain: $\beta = 0.001$ is the recommended default across all datasets, and the uncertainty AUROC is roughly insensitive to β , so the UQ signal is robust to this choice even when the alignment metrics themselves shift. End-to-end wall time on a single NVIDIA RTX 3080 with 16 GB RAM is ~ 120 s per PBMC run, ~ 50 s per Brain run, and ~ 95 s per CITE-seq run; the Sinkhorn step uses a seeded 5000-cell subsample for the uncertainty analysis (full-dataset metrics FOSCTTM/LT/ARI are computed without subsampling), and memory peak is ~ 8 GB at 5000×5000 cost matrix in float64. All seeds are set deterministically (NumPy, PyTorch, cuDNN); multi-seed results (mean $\pm$ std) are aggregated over seeds 0–9 on PBMC $\beta = 0.001$ and seeds 0–2 on all other configurations; the uniPort baseline is run from a separate Python virtual environment with `numpy<2` and `anndata 0.11.4` because uniPort 1.3 depends on the deprecated `np.Inf`, and code, configs, and all aggregate JSON result files will be linked upon de-anonymization.




Figure 12: β sensitivity on PBMC and Brain. $\beta = 0.001$ is the recommended default; the UQ AUROC is roughly insensitive to β , so the UQ signal is robust to this choice.